

| *Is there real demand? A studio, or another AI playground?*

A studio market, repricing quarterly

- **\$5B** in 2023 → **\$18.6B** in 2026.
- Two **\$5B+** companies minted in one year.
- Enterprise is now the **majority** of Higgsfield revenue.
- Agents next: **42×** agentic usage in 3 months.

7 MONTHS

\$1.3B → \$5.4B

Higgsfield valuation

12 MONTHS

\$20M → \$700M

annualized revenue

FEB 2026

\$5.3B

Runway valuation

FORTUNE 500

390

use Higgsfield

What is graph-based inference?

A **graph** is a production line. A **model** is a part on the shelf.


IF YOU KNOW VLLM

There, the model is a **deployment**. Here, it is a **part** pulled from the shelf per job.

IF YOU DON'T

A prompt box is a **vending machine**. A graph is a **production line**: change one station, only that one reruns.

- One job: 10+ images, text, LoRAs, 2–3 models, 40 nodes.
- Most nodes are not inference.
- Templates: built by others, or yourself. The shelf only grows.



- 1 template
- 2 model loader
- 3 missing models
- 4 Run: the GPU step

| *Why ComfyUI, against the commercial tools?*

The open-source leader, missing only its enterprise story

- **4M** users · **150k** downloads/day · **\$500M** valuation.
- In production at Amazon Studios, Apple, Netflix, Nike, Ubisoft.
- Open source under a price umbrella: **the house strategy**.
- “ComfyUI artist” is a job title.

	HIGGSFIELD · KREA · RUNWAY · MIDJOURNEY	COMFYUI ON THIS POOL
Price	\$79–215 per seat per month	Open source, on your GPUs
Models	The vendor's catalog	Any model, yours or public
Nodes	The vendor's pipeline	60,000+ community nodes
Your data	The vendor's cloud	Your VPC, your SSO

| Can the machine under each desk be broken up?

A card per person, **idle most of the day**

- \$6–8k desktop per seat, or a cloud card each: **~\$7,100/mo** for ten.
- One shared box: one unauthenticated process.
- NDA studios cannot use a vendor cloud.

PER SEAT

\$6–8k

a desktop

TEN, CLOUD

~\$7,100/mo

a card each

Load a template



See missing models



Load the models



Wire & fix the graph




Run

Four of five steps never touch a GPU.

GPUs to the datacenter, frontend untouched?

Same application. **Only the silicon moved.**


- Stock ComfyUI, pinned. Same canvas, same loop.
- Models pulled once: one shelf for the whole floor.
- Workers scale 0 → N on queue depth.



| What drives GPU cost?

Pooling is **queueing**, not division

- Ten designers: **~\$7,100/mo** as cards, **\$120–950** as a pool.
- Zero **nodes**, not zero pods, is what saves money.



SAVINGS AGAINST A CARD EACH, WARM 9–6

5 PEOPLE

33%

10 PEOPLE

45%

30 PEOPLE


64%

Doubling the team does not double the overlap.

Where it loses: cold start without a warm floor, custom nodes need a rebuild.
README, "Where this loses".

Why not an engine per model?

Residency scales with the **catalog**. The pool scales with **concurrency**.




- Every model on an engine: **+\$713/mo** each.
- The pool: **flat**.
- The marginal model: a machine there, a disk here.

What does the repo add beyond the Omni bridge?

The production layer: **queue, isolation, storage, scale**


- **WebSocket** in, **Redis queue** out. No browser touches a worker.
- No shared storage? **S3 + node cache**. Hundreds of designers, same shape.
- **Shared storage**: a model pulled once is there for everyone.
- **Workers unreachable** by design. SSO, audit, showback, quota.



Where does vLLM Omni fit? Does this touch the application layer?

Two tiers, one stack. **Showback decides** when a model moves.

- Engine = **throughput** tier. Pool = **flexibility** tier.
- One model dominates showback → give it an **engine**.
- Omni is **one more node** on the canvas. The bridge exists today.
- ComfyUI stays upstream's app. This repo is **wiring**.



Why is this ours to do?

Red Hat is the best place to run ComfyUI

- An **\$18B** category with no enterprise deployment story.
- OpenShift already had the parts: GPU Operator, SSO, scale-to-zero pools, network policy, audit.
- Same app, real workflows, constrained GPUs, cost-effective.
- Inference needs a **graph-input** answer beside the text one.

ComfyUI

upstream, unchanged

The pool – this repo

queue · isolation · storage · showback

OpenShift

SSO · machine pools · network policy · audit

NVIDIA GPU Operator

L4 → L40S → H100

The graph-input half of inference, on the platform we already sell.