



What is ComfyUI, and why is it hard to share?

- The open-source standard for AI images and video. Netflix, Apple, Nike use it.
- Paid tools: **\$79–215 a seat**, your data in their cloud.
- A desktop app: **no login**, add-ons **run any code**.
- So teams give everyone a GPU, idle most of the day.



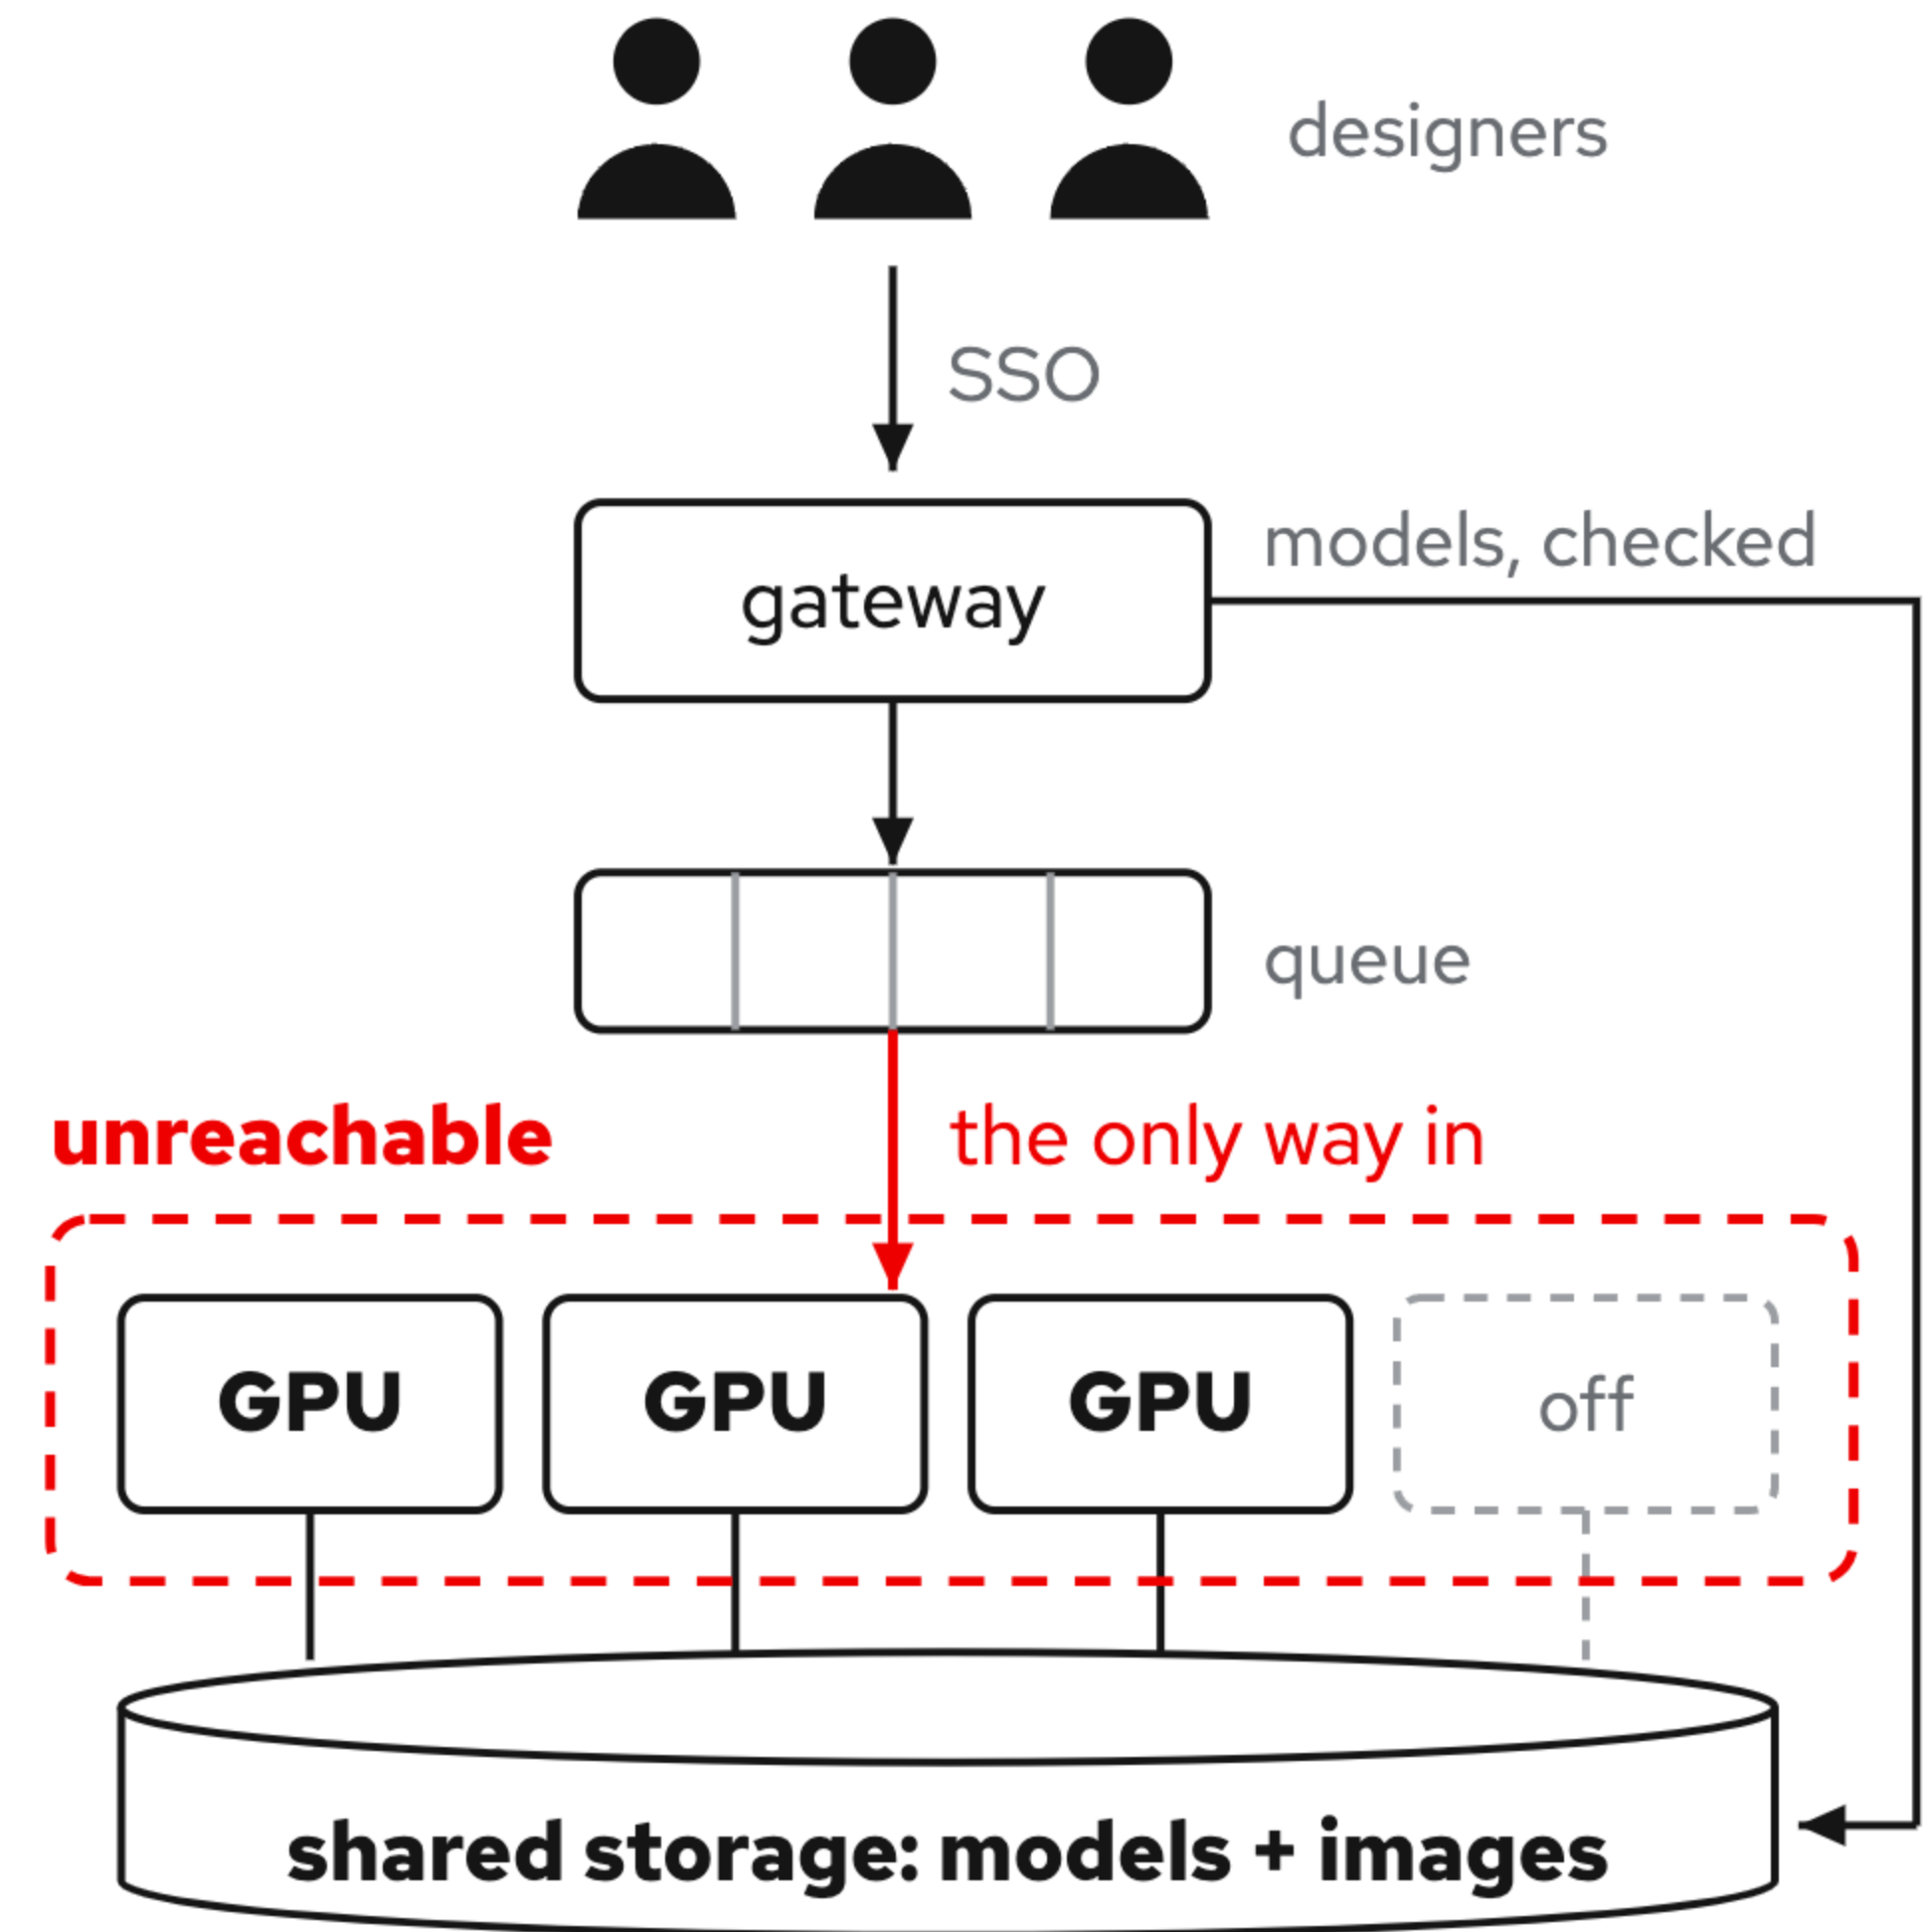
So what's the design?

- Same ComfyUI, same canvas. **Only the GPUs move.**
- Every Run becomes a job on one queue.
- A pool of GPUs takes the jobs. Idle? **Zero machines.**
- Models download once, to shared storage.



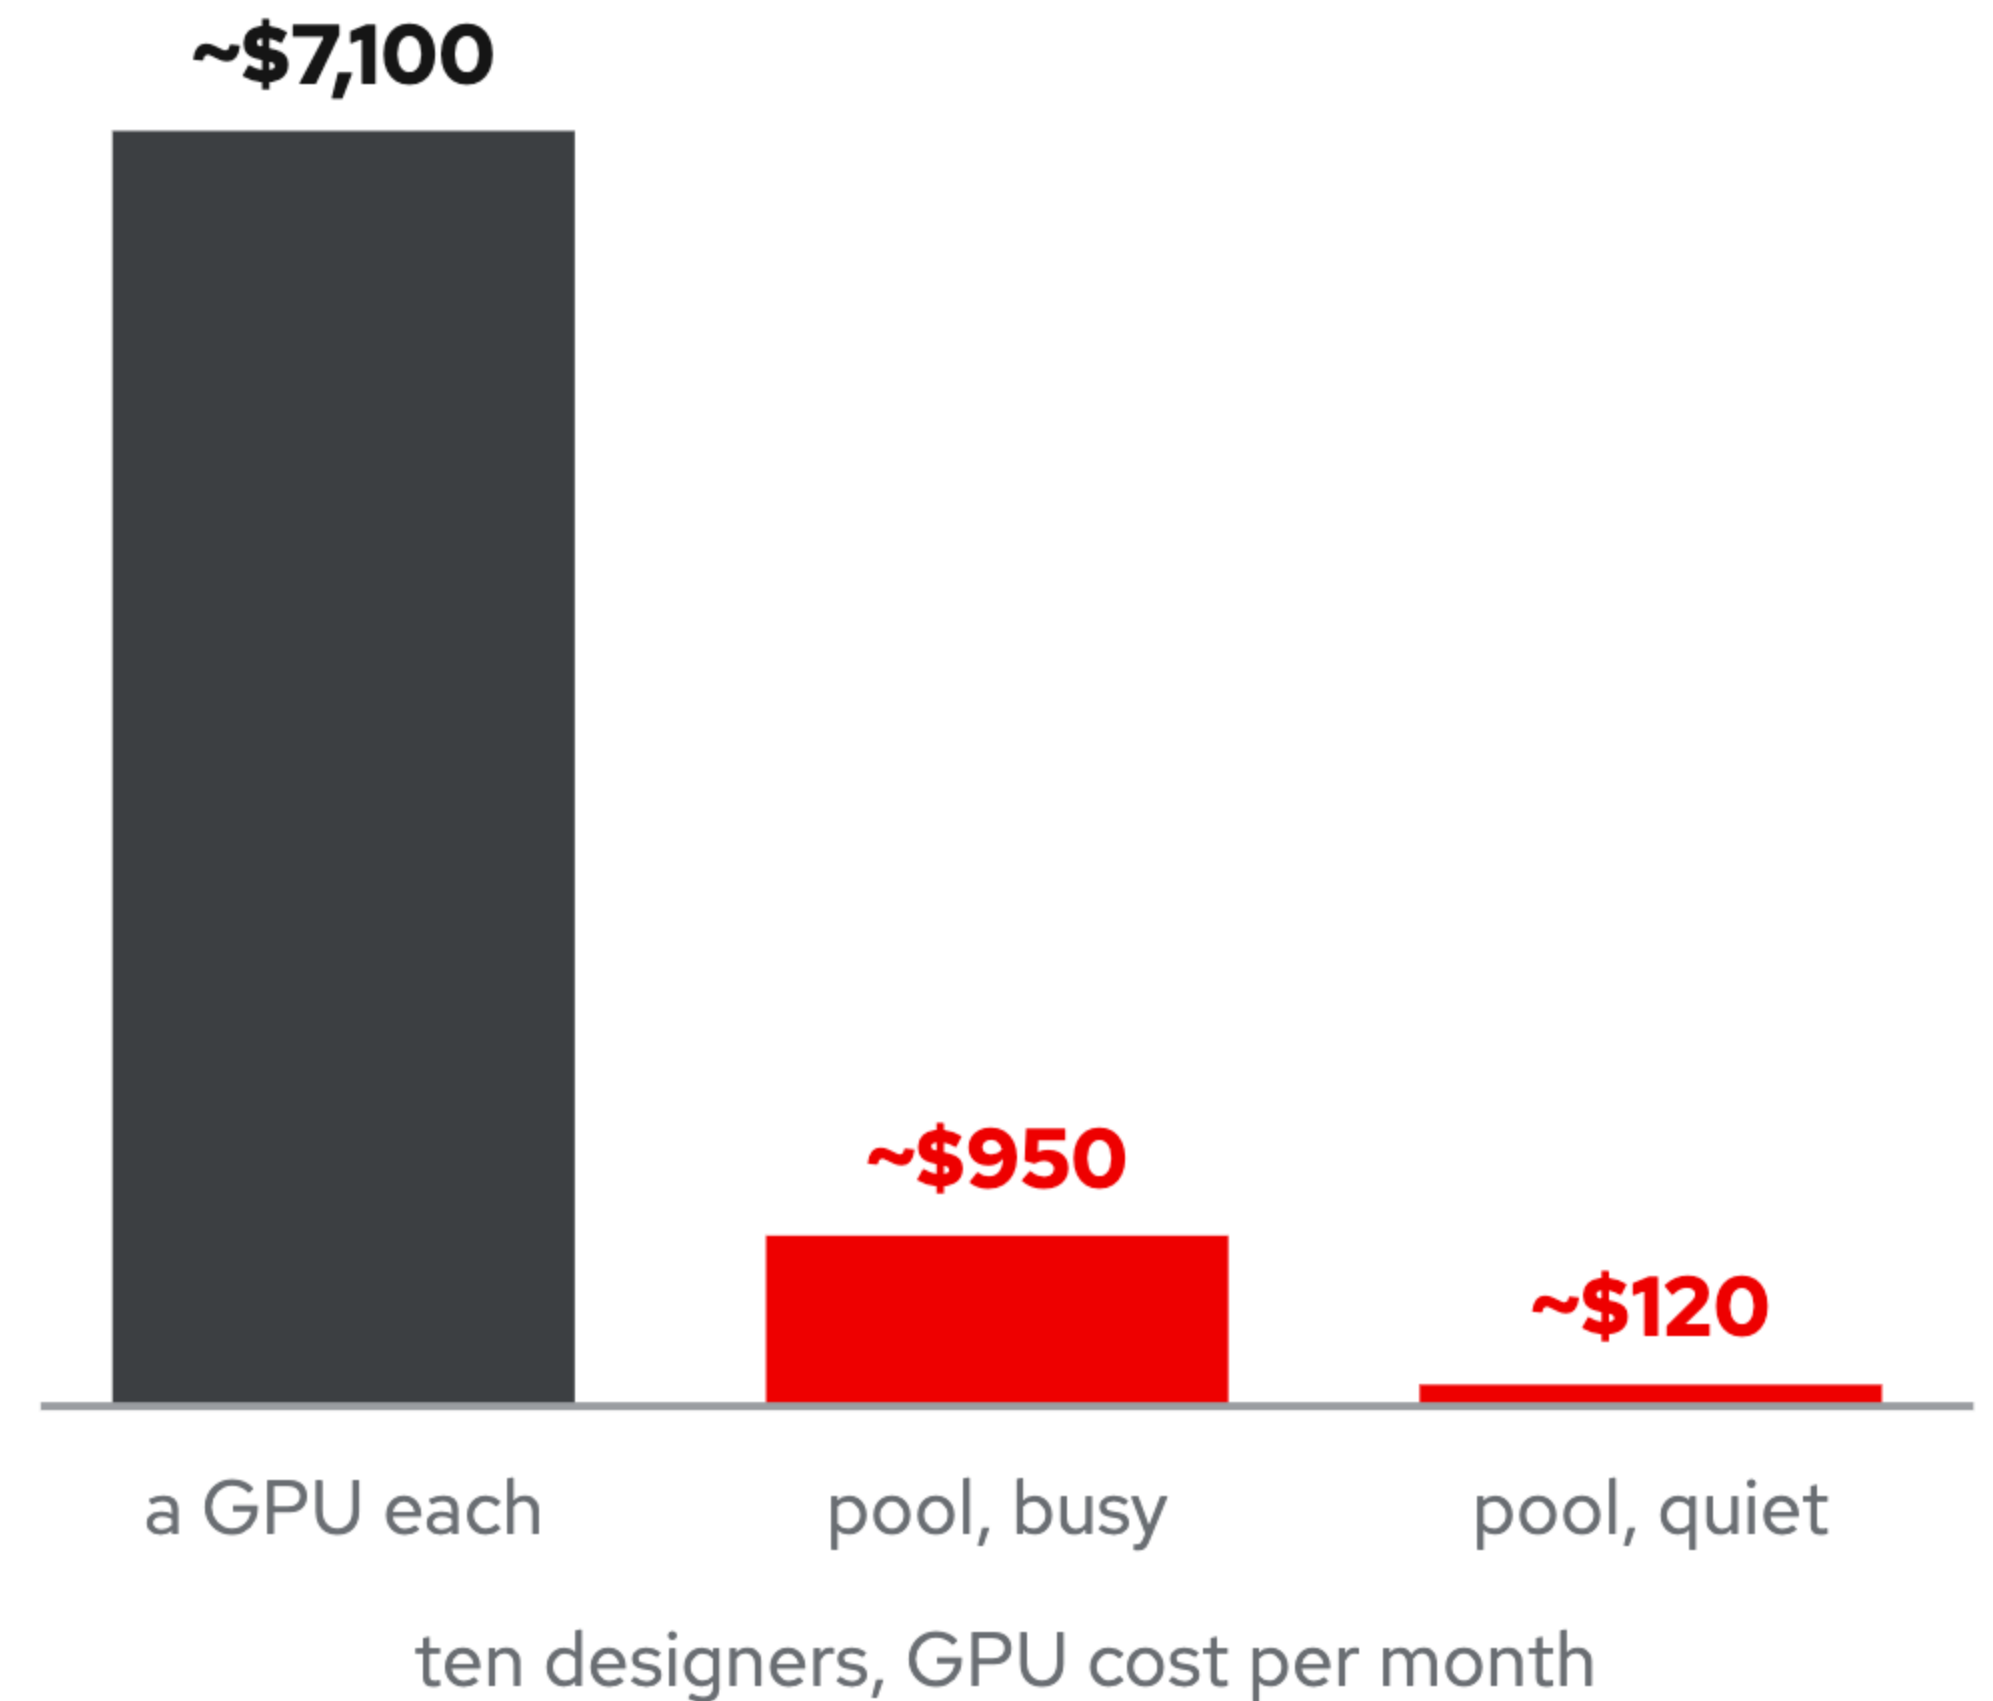
ComfyUI runs any code. How is sharing safe?

- GPU workers are **unreachable**, not defended.
- No network address. The only way in is the queue.
- No internet out: models come in through the gateway, checked.
- Company SSO at the door. GPU time metered per person.



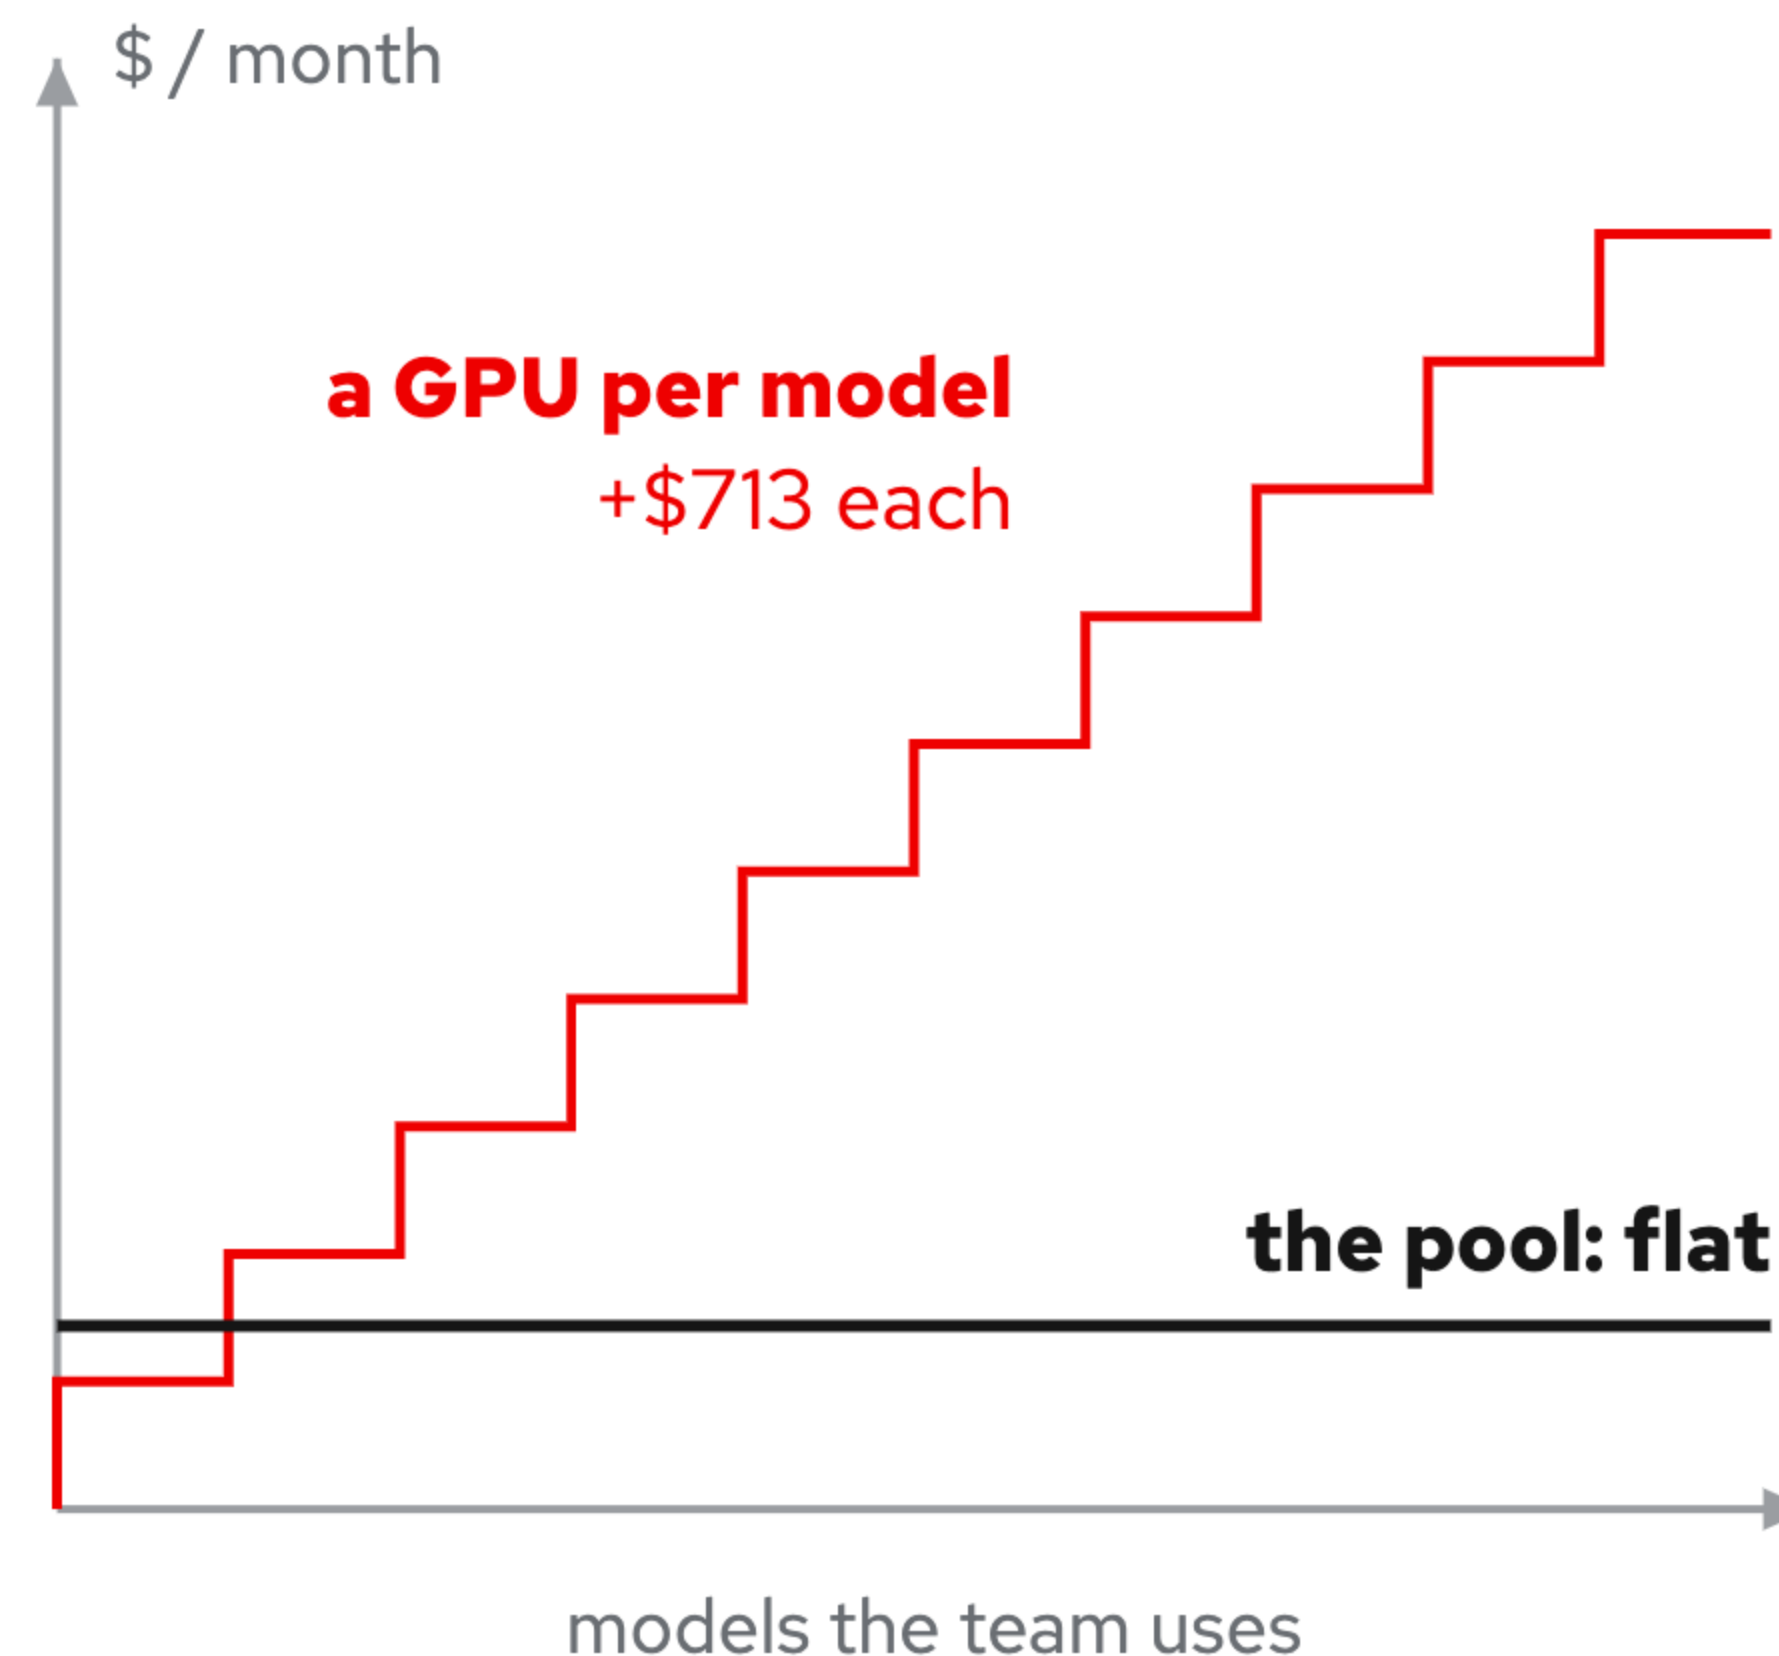
What does it cost, and what's the catch?

- Ten designers: **~\$7,100/mo** as a GPU each, **\$120–950** as a pool.
- Works because few people hit Run at once.
- Catch: 8–17 min cold start. One warm GPU fixes it, ~\$190/mo.
- Catch: new add-on code means an image rebuild.



Why not serve each model on vLLM?

- vLLM keeps a model loaded on a GPU. Best for **one busy model**.
- A design job is a graph: 2–3 models, many non-model steps.
- Each extra model on vLLM: **+\$713/mo**. In the pool: disk.
- Use both: move a model to vLLM Omni once it dominates usage.



Where is it, and what do we need?

- ☒ Built: ~9,000 lines of wiring. OpenShift already had the parts.
 - ☒ Ran end to end on real GPUs. Also runs on a laptop.
 - ☐ Not proven at scale: a handful of users, three workers.
- Ask: **an owner**, a 30-minute PM read, **~\$50** for a cluster day.